

AI Coding with Junie

JetBrains' Coding Agent for IDE, CLI, and CI

Plan, implement, review, and verify
with a coding agent that understands your project

What You'll Master Today

Core Concepts

-  **Ask Mode:** Read-only exploration
-  **Plan Mode:** Align before coding
-  **Code Mode:** Implement and verify
-  **Safety:** Approvals, allowlists, review

Hands-On Labs

-  **Java:** Spring Boot APIs
-  **Python:** Refactoring & testing
-  **React:** Forms & validation
-  **MCP:** Current docs and tools

Prerequisites Check

- ✓ **JetBrains IDE** (IntelliJ IDEA, PyCharm, WebStorm, GoLand, PhpStorm, RustRover, RubyMine)
- ✓ **JetBrains AI** subscription, Junie API key, or BYOK provider key
- ✓ **Development Environment** (Java 21+ / Python 3.8+ / Node 18+)

Cost control: Use `/usage``, choose models intentionally, and prefer BYOK/local models for heavy practice

💡 **Quick Setup:** We'll verify Junie, CLI, and lab runtimes together at the start

Getting Started — Two Surfaces

JetBrains IDE

1. AI Chat → Agent dropdown → ****Junie****
2. Separate Junie tool window remains available
3. Best for indexes, refactorings, debugger, inspections

Junie CLI

1. ``brew tap jetbrains-junie/junie && brew install junie``
2. Or: ``curl -fsSL https://junie.jetbrains.com/install.sh | bash``
3. Run ``junie`` in any project, or ``junie "task"`` headlessly

Model and provider options:

JetBrains account / credits

OpenAI / Anthropic / Google / xAI

OpenRouter / Copilot


Ollama / LM Studio / LiteLLM

Use ``/model`` and ``/usage`` to stay intentional



What is Junie?

A coding agent that works where JetBrains developers already work



NEW

Junie After Beta

General Availability Story

- JetBrains now calls Junie a coding agent
- IDE, terminal, and CI/CD are one product story
- Plans before it codes
- Runs long tasks while you review intent

Plan Mode Is Central

- Read-only exploration first
- Requirements, design, tests, delivery steps
- Review and refine before implementation
- Save plans as durable task docs


Shared Agent Context

- ``.junie/AGENTS.md`` first
- Root ``AGENTS.md`` also supported
- Skills live in ``.junie/skills/``
- Extensions bundle skills, MCP, commands, agents

Beyond the IDE

- ``.review`` for local diffs
- ``.debug`` with JetBrains debugger context
- ``.remote`` to continue a CLI session in browser
- GitHub Action and GitLab CI workflows

Junie Surfaces



The CLI can connect to a matching JetBrains IDE for symbol-aware search, inspections, test discovery, and debugger-backed workflows.

Three Working Modes



Ask Mode

Read-Only Analysis

- Explain complex code
- Analyze architecture
- Find bugs & issues
- Review security
- Check test coverage



Plan Mode

Design Before Edits

- Requirements
- Technical design
- Testing strategy
- Delivery steps
- Confirm or refine




Code Mode

Make Changes

- Implement features
- Fix bugs
- Refactor code
- Generate tests
- Update dependencies

💡 Professional Workflow



Use Plan mode when the cost of the wrong implementation is higher than the cost of five minutes of alignment.



Plan Mode in Practice

Start It

- Press `Shift+Tab` in Junie CLI
- Or type `/plan <task>`
- In the IDE, start with Ask/Auto and request a plan first
- For installed flags, check `junie --help`

Review It

- Read assumptions
- Edit scope before code exists
- Check testing strategy
- Confirm implementation only after the plan matches your intent

Teaching move: use one complete prompt, let Junie plan, then ask attendees what they would change before implementation.



Project Guidelines

Preferred: `.junie/AGENTS.md` | **Also supported:** `root AGENTS.md` | **Legacy:**
`.junie/guidelines.md`

```
## Technology Stack
- Framework: Spring Boot 3.2
- Testing: JUnit 5 + AssertJ

## Conventions
- REST: /api/v1/{resource}
- DTOs: Java records
- Tests: Given-When-Then
```

Lookup order: `.junie/AGENTS.md` → `AGENTS.md` (project root) → `.junie/guidelines.md` (legacy). Global guidelines can live in `~/junie/AGENTS.md`.

✗ Without Guidelines

Inconsistent patterns

Multiple review cycles

Style conflicts


✓ With Guidelines

Consistent output

70% fewer reviews

Team alignment

Safety Controls



⚡ When to Use Each Mode

Mode	Good For	Not For
Approvals	• Critical business logic • First-time tasks • Learning Junie	• Repetitive formatting • Simple refactors
Allowlist	• Known file sets • Defined scope • Team reviews	• Exploratory changes • Unknown impact
Brave	• Test generation • Formatting • Documentation	• Production code • Database changes • First attempts

MCP: Model Context Protocol

Use MCP For

- External services
- Current docs
- Browser exploration
- Shared project tools

Use Skills For

- Repeatable workflows
- Test conventions
- Review checklists
- Team patterns

```
{
  "mcpServers": {
    "context7": {
      "command": "npx",
      "args": ["-y", "@upstash/context7-mcp"]
    },
    "playwright": {
      "command": "npx",
      "args": ["-y", "@playwright/mcp"]
    }
  }
}
```

Where it lives


CLI: `~/mcp``

Project: `~/junie/mcp/mcp.json``

User: `~/junie/mcp/mcp.json``

context7 is the main workshop MCP example.

Lab Overview



Java

Spring Boot

REST APIs

JUnit + AssertJ



Python

PEP 8

Type hints

pytest



React

TypeScript

Forms

Testing Library



MCP

context7

Upgrades

Migration



Skills

Reusable

Workflow

Playwright



Lab E

Weather API

Proxy

Optional

Lab A: Spring Boot Journey



Baseline Pass

- Ask Junie to inspect first
- Let it propose a plan
- Watch where style drifts

Guided Pass

- ``.junie/AGENTS.md`` loaded
- One complete feature prompt
- Tests and review included

Lab B: Python Transformation

Starting Point

```
def calc(x,y,op):  
    if op=="add": return x+y  
    elif op=="sub": return x-y
```

- No type hints
- Poor naming
- No tests
- No docs

Target Outcome

```
def calculate(  
    left: float,  
    right: float,  
    operation: Operation,  
) -> float:  
    """Apply a supported operation."""
```

- Type-safe API
- pytest coverage
- Documented edge cases



Lab C: React Forms

```
interface RegistrationForm {  
  email: string;    // valid email  
  password: string; // min 8, special char  
  confirm: string;  // must match  
  terms: boolean;   // required  
}
```

React Hook Form
Form state


Zod
Validation

Testing Library
90% coverage

Accessibility
ARIA compliant

Teaching move: Use one complete prompt that asks for the component, validation, accessible errors, tests, and a final self-review.

Lab D: Smart Upgrades with context7



React 17 → React 18
Legacy Modern

Four-hour version: emphasize MCP setup, migration plan quality, and one contained implementation slice.

Mini-Lab: Create a Playwright Skill

Capture the workflow once, reuse it every time

Create a skill named playwright-e2e that captures our browser testing conventions:

- Prefer accessible locators
- Use page objects only after duplication appears
- Run `npx playwright test`
- Summarize failures and screenshots

1. Inspect Existing Skills

2. Ask Junie to Create One

3. Use It on a Test

4. Run Local CLI

Teaching point: Playwright MCP is good for interactive browser exploration. A Skill is better for repeatable team testing conventions.

Common Patterns

TDD Flow

1. Ask: "What should fail?"
2. Plan: "How will we prove it?"
3. Code: "Implement and run tests"

Red → Green → Refactor

Plan-First Feature

1. State acceptance criteria
2. Review Junie's plan
3. Approve a focused diff

Scope stays visible

Upgrades

1. Use current docs
2. Plan migration
3. Verify behavior

Systematic migration

CI/CD Integration — Junie GitHub Action

```
name: Junie Agent
on:
  issue_comment:
    types: [created]
  pull_request_review_comment:
    types: [created]

permissions: { contents: write, pull-requests: write, issues: write }

jobs:
  junie:
    if: contains(github.event.comment.body, '@junie-agent')
    runs-on: ubuntu-latest
    steps:
      - uses: JetBrains/junie-github-action@v0
        with:
          junie_api_key: ${ secrets.JUNIE_API_KEY }
```

GitHub: `@junie-agent` in issues or PRs

GitLab: `#junie` in merge request comments

CLI: `junie --task "prompt"` for headless local automation



Remote Mode Demo

What It Is

- Browser UI for the same CLI session
- Works from laptop, tablet, or phone browser
- Terminal and web UI stay synchronized
- Your machine must stay awake

Demo Flow

1. Start ``junie`` in the terminal
2. Run ``/remote``
3. Open ``junie.jetbrains.com/remote``
4. Respond from the browser
5. Run ``/remote`` again to stop

Remote mode uses the Junie service. JetBrains Account sign-in is the simplest path; a Junie API key is the alternate route.

🤔 Why Junie Instead of Cursor?

Use Junie When...

- Team lives in JetBrains IDEs
- IDE refactorings matter
- Plans before edits matter
- IDE, CLI, and CI should align

Use Cursor When...

- AI-first editor is desired
- Cursor Tab is central
- Background agents matter
- Team already uses rules/Bugbot

Compare Live


- Same repo
- Same task
- Same acceptance criteria
- Stop before a second workshop

Cursor can also run inside JetBrains through ACP, so this is about workflow fit.

✨ Key Takeaways

- ✓ **Ask before Code** - Understand first, implement second
- ✓ **Plan before big changes** - Align on requirements, tests, and scope
- ✓ **Guidelines drive consistency** - Encode your team's DNA
- ✓ **Skills preserve workflows** - Reuse task-specific patterns
- ✓ **Build trust gradually** - Approvals → Allowlist → Brave
- ✓ **CLI matters** - Use Junie outside the IDE and in automation
- ✓ **Safety first** - Review diffs, run tests, maintain control

Your Next Steps



 Start with Ask and Plan → Approve focused diffs → Add guidelines → Automate later

Discussion Time

Let's Explore Together:

-  IDE plugin vs CLI vs CI/CD — which fits your workflow?
-  Where do you always want human review?
-  How could AGENTS.md guidelines help your team?
-  Which MCP tools and Agent Skills would you create?
-  Where does Cursor, Codex, Claude, or Antigravity fit better?

Resources Hub

Documentation

-  [Junie Docs](#)
-  [Guidelines & Memory](#)
-  [Agent Skills](#)
-  [CLI Reference](#)

Agent Workflows

-  [GitHub Action](#)
-  [GitLab CI](#)
-  [Plan Mode](#)
-  [Remote Mode](#)

GitHub

-  [Guidelines Catalog](#)
-  [Context7 MCP](#)
-  [This repository](#)

Adjacent Tools

-  [Cursor Docs](#)
-  [Codex / Claude / Antigravity docs](#)
-  [Agent Communication Protocol](#)

Community

-  [JetBrains AI Slack](#)
-  [Stack Overflow: `jetbrains-junie`](#)
-  [Support team](#)

 Thank You!

AI amplifies good practices

 Remember:

**Good agents amplify clear intent,
tested code, and team conventions.**

Questions? Let's explore together! 

About Ken Kousen

Ken Kousen

President, Kousen IT, Inc.

 ken.kousen@kousenit.com

 github.com/kousen

 [@talesfromthejarside](https://twitter.com/talesfromthejarside)

 kousenit.substack.com

 linkedin.com/in/kenkousen

 bsky.app/profile/kousenit.com

AI agents aren't here to replace developers—
they're here to make us **better** developers.

Bonus: Quick Reference

IDE Shortcuts

- ``Ctrl+Alt+J`` - Open Junie
- ``Ctrl+Enter`` - Execute
- ``Esc`` - Cancel operation ### CLI

Shortcuts • ``Shift+Tab`` - Plan mode

- ``/mcp`` - MCP setup
- ``/model`` - Choose model
- ``/usage`` - Credits usage
- ``Ctrl+B`` - Brave mode
- ``Ctrl+R`` - Prompt history

Common Commands

"Analyze this file"

"Create a plan first"

"Add tests with 90% coverage"

"Implement and run tests"

"Review the diff"

MCP Tools

```
{  
  "context7": "Library docs",  
  "playwright": "E2E tests",  
  "search": "Web search",  
  "custom": "Your tools"  
}
```

 Print this slide for your desk!